CS 5/7322 Spring 2025

Presentations

All students registered for CS 7322 are required to present a paper related to the project (in addition to the final project presentation).

Presentations

The presentation will be based on project groups. Each student in the project group will present a paper (on occasion, two people will present one paper) that is mostly in the list of papers to be read. Each group will decide among its members who to present which paper.

The presentation should be about 15 minutes with about 2 extra minutes for questions.

For those who are distance students, you should record a 15-minute presentation of the paper and upload it to Canvas (under the "Presentation" tab). You can also upload it to a site (e.g. YouTube) and provide a link to it. You should get it done 24 hours before your presentation date so I can have a check before it.

Each person should upload the slide of the presentation at least 24 hours before the presentation also.

Attendance for CS 5322 students

All students taking 5322 will have to be present. Roll call will be taken, and the number of lectures that you are here will count towards your presentation grade.

Papers and presentation schedules

- Chen/Yu (Project 2) {Date: 4/21 (Mon)]
    - [Two students combined to present this (30 minutes total)] Pengfei Liu, Weizhe Yuan, Jinlan Fu, Zhengbao Jiang, Hiroaki Hayashi, and Graham Neubig. 2023. Pre-train, Prompt, and Predict: A Systematic Survey of Prompting Methods in Natural Language Processing. ACM Comput. Surv. 55, 9, Article 195 (September 2023), 35 pages. https://doi.org/10.1145/3560815

- Bryan/Shellet/Zhao (Project 4) [Date: 4/21 (Mon) – 2 papers, 4/23 (Wed) – 1 paper]
    - Rongwu Xu, Zehan Qi, Zhijiang Guo, Cunxiang Wang, Hongru Wang, Yue Zhang, and Wei Xu. 2024. Knowledge Conflicts for LLMs: A Survey. In Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing, pages 8541–8565, Miami, Florida, USA. Association for Computational Linguistics.
    - Hou, Yufang, et al. "WikiContradict: A Benchmark for Evaluating LLMs on Real-World Knowledge Conflicts from Wikipedia." *Advances in Neural Information Processing Systems* 37 (2025): 109701-109747.
    - Xie, Jian, et al. "Adaptive chameleon or stubborn sloth: Revealing the behavior of large language models in knowledge conflicts." *The Twelfth International Conference on Learning Representations*. 2023.

- Allen/McCullough/Miller (Project 9) [Date 4/23 (Wed)]
    - Alturayeif, N., Luqman, H. & Ahmed, M. A systematic review of machine learning techniques for stance detection and its applications. *Neural Comput & Applic* **35**, 5113–5144 (2023).
    - Mayor, E., Miani, A. A topic models analysis of the news coverage of the Omicron variant in the United Kingdom press. *BMC Public Health* **23**, 1509 (2023).
    - Gül, Ilker, Rémi Lebret, and Karl Aberer. "Stance detection on social media with fine-tuned large language models." *arXiv preprint arXiv:2404.12171* (2024).

- Al-Atoum/Carpenter/Chaudhari (Project 8) [Date 4/28 (Mon)]
    - G. Koutrika, A. Simitsis and Y. E. Ioannidis, "Explaining structured queries in natural language," *2010 IEEE 26th International Conference on Data Engineering (ICDE 2010)*, Long Beach, CA, USA, 2010, pp. 333-344,
    - Kun Xu, Lingfei Wu, Zhiguo Wang, Yansong Feng, and Vadim Sheinin. 2018. SQL-to-Text Generation with Graph-to-Sequence Model. In *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing*, pages 931–936, Brussels, Belgium. Association for Computational Linguistics.
    - V. Câmara, R. Mendonca-Neto, A. Silva and L. Cordovil, "A Large Language Model approach to SQL-to-Text Generation," *2024 IEEE International Conference on Consumer Electronics (ICCE)*, Las Vegas, NV, USA, 2024, pp. 1-4, doi: 10.1109/ICCE59016.2024.10444148.
    -
- Alam/Mejia/Miller (Project 3) [Date: 4/30 (Wed)]
    - [Two students combined to present this (30 minutes total)] Gao, Yunfan, et al. "Retrieval-augmented generation for large language models: A survey." *arXiv preprint arXiv:2312.10997* 2 (2023). (2 student combine to present the paper: total 25-30 minutes)
    - Du, Mingzhe, et al. "From static to dynamic: A continual learning framework for large language models." arXiv preprint arXiv:2310.14248 (2023).

- Vats [Date 4/30 (Wed)]
    - TBD